

More layers

CONVOLUTIONAL LAYERS



layer_conv_lstm_2d()

Convolutional LSTM

layer_separable_conv_2d()

Depthwise separable 2D

layer_upsampling_1d()

layer_upsampling_2d()

layer_upsampling_3d()

Upsampling layer

layer_zero_padding_1d()

layer_zero_padding_2d()

layer_zero_padding_3d()

Zero-padding layer

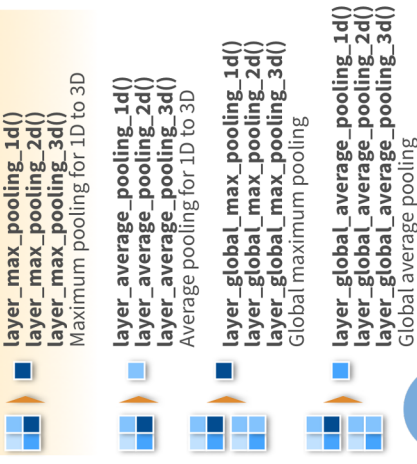
layer_cropping_1d()

layer_cropping_2d()

layer_cropping_3d()

Cropping layer

POOLING LAYERS



Preprocessing

SEQUENCE PREPROCESSING

pad_sequences()
Pads each sequence to the same length (length of the longest sequence)

skipgrams()

Generates skipgram word pairs

make_sampling_table()

Generates word rank-based probabilistic sampling table

TEXT PREPROCESSING

text_tokenizer() Text tokenization utility

fit_text_tokenizer() Update tokenizer internal vocabulary

save_text_tokenizer(); load_text_tokenizer()

Save a text tokenizer to an external file

texts_to_sequences();

texts_to_sequences_generator()

Transforms each text in texts to sequence of integers

texts_to_matrix(); sequences_to_matrix()

Convert a list of sequences into a matrix

text_one_hot() One-hot encode text to word indices

text_hashing_trick()

Converts a text to a sequence of indexes in a fixed-size hashing space

text_to_word_sequence()

Convert text to a sequence of words (or tokens)

IMAGE PREPROCESSING

image_load() Loads an image into PIL format.

flow_images_from_data()

flow_images_from_directory()
Generates batches of augmented/normalized data from images and labels, or a directory

image_data_generator() Generate minibatches of image data with real-time data augmentation.

fit_image_data_generator() Fit image data generator internal statistics to some sample data

generator_next() Retrieve the next item

image_to_array(); image_array_resize()

image_array_save() 3D array representation



Pre-trained models

Keras applications are deep learning models that are made available alongside pre-trained weights. These models can be used for prediction, feature extraction, and fine-tuning.

application_xception()

xception_preprocess_input()

Xception v1 model

application_inception_v3()

inception_v3_preprocess_input()

Inception v3 model, with weights pre-trained on ImageNet

application_inception_resnet_v2()

inception_resnet_v2_preprocess_input()

Inception-ResNet v2 model, with weights trained on ImageNet

application_vgg16(); application_vgg19()

VGG16 and VGG19 models

application_resnet50() ResNet50 model

application_mobilenet()

mobilenet_preprocess_input()

mobilenet_decode_predictions()

mobilenet_load_model_hdf5()

MobileNet model architecture

IMAGENET

ImageNet is a large database of images with labels, extensively used for deep learning

imagenet_preprocess_input()

imagenet_decode_predictions()

Preprocesses a tensor encoding a batch of images for ImageNet, and decodes predictions

Callbacks

A callback is a set of functions to be applied at given stages of the training procedure. You can use callbacks to get a view on internal states and statistics of the model during training.

callback_early_stopping() Stop training when a monitored quantity has stopped improving

callback_learning_rate_scheduler() Learning rate scheduler

callback_tensorboard() TensorBoard basic visualizations